G.2 Mistral-7B

In this section, we present the results for the Mistral-7B-Instruct-v0.3 model. As shown in figure

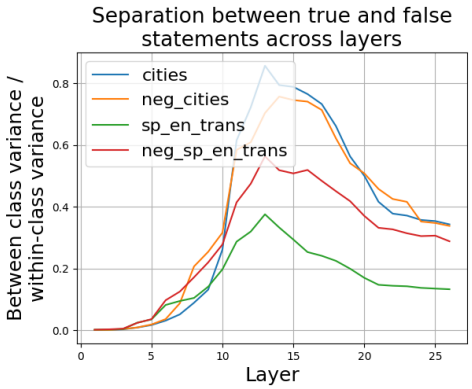

Figure 16: Mistral-7B: Ratio between the between-class variance and within-class variance of activations corresponding to true and false statements, across residual stream layers.

16, the largest separation between true and false statements occurs in layer 13. Therefore, we use activations from layer 13 for the subsequent analysis of the Mistral-7B model.

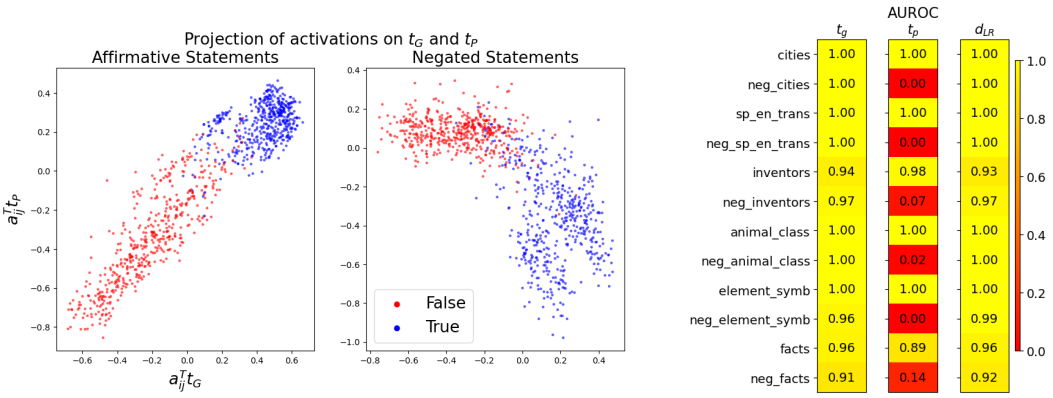

(a) (b)

Figure 17: Mistral-7B: Left (a): Activations $\mathbf{a}_{ij}$ projected onto $\mathbf{t}_G$ and $\mathbf{t}_P$. Right (b): Separation of true and false statements along different truth directions as measured by the AUROC, averaged over 10 training runs.