G.4 Gemma-2-27B

In this section, we present the results for the Gemma-2-27B-Instruct model. As shown in figure 28,

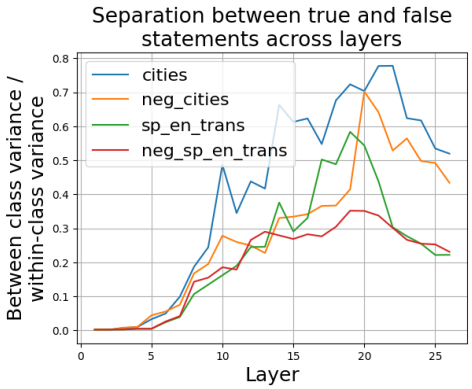

Figure 28: Gemma-2-27B: Ratio between the between-class variance and within-class variance of activations corresponding to true and false statements, across residual stream layers.

the largest separation between true and false statements occurs approximately in layer 20. Therefore, we use activations from layer 20 for the subsequent analysis of the Gemma-2-27B-Instruct model.

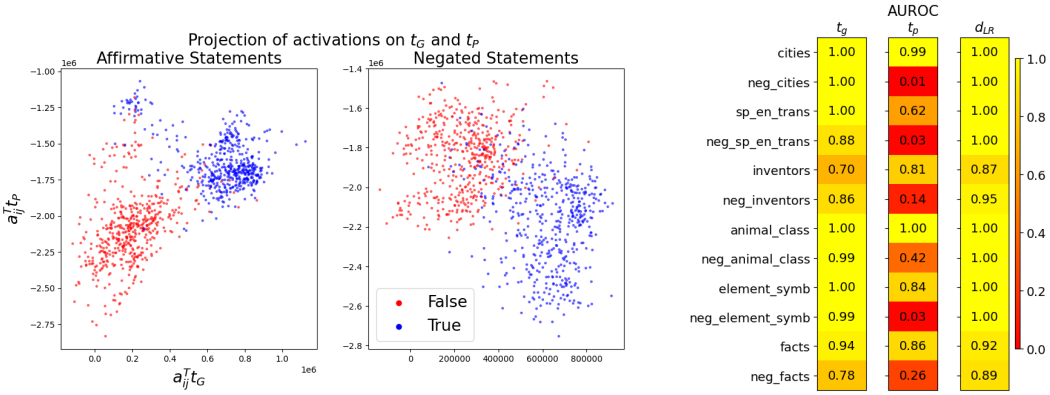

(a) (b)

Figure 29: Gemma-2-27B: Left (a): Activations $\mathbf{a}_{ij}$ projected onto $\mathbf{t}_G$ and $\mathbf{t}_P$. Right (b): Separation of true and false statements along different truth directions as measured by the AUROC, averaged over 10 training runs.