G.5 LLaMA3-8B-base

In this section, we present the results for the LLaMA3-8B base model. As shown in figure 34, the

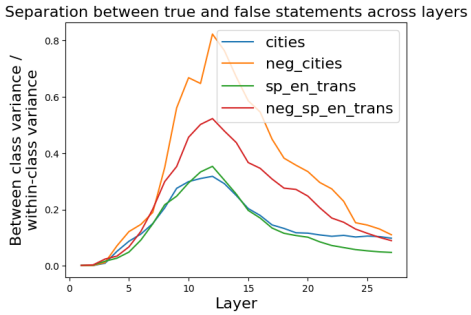

Figure 34: LLaMA3-8B-base: Ratio between the between-class variance and within-class variance of activations corresponding to true and false statements, across residual stream layers.

largest separation between true and false statements occurs in layer 12. Therefore, we use activations from layer 12 for the subsequent analysis of the LLaMA3-8B-base model.

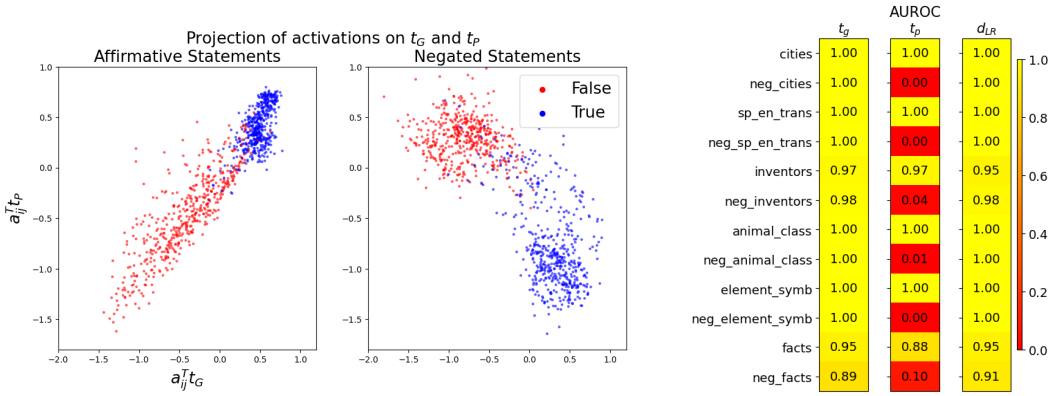

(a) (b)

Figure 35: LLaMA3-8B-base: Left (a): Activations $\mathbf{a}_{ij}$ projected onto $\mathbf{t}_G$ and $\mathbf{t}_P$. Right (b): Separation of true and false statements along different truth directions as measured by the AUROC, averaged over 10 training runs.